

CMBLensing.py: A Python package for MAP Estimation of CMB Fields and Cosmological Parameter Sampling

Zane Blood^{1*}, Kimmy Wu^{1*}, and Fei Ge^{1*}

¹ California Institute of Technology * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))

Summary

CMBLensing.py (Blood, 2026) is the Python migration of **CMBLensing.jl** (Millea, 2017), a Julia package for simulating Cosmic Microwave Background (CMB) fields, gravitational lensing, Maximum A Posteriori (MAP) estimates of CMB fields, and cosmological parameter sampling. Written in JAX (Bradbury et al., 2018), most of its methods are Just-In-Time (JIT) compilable and run on CPU or GPU without special configuration. **CMBLensing.py** reproduces the full functionality of the original code and has been extensively tested for regressions. We have also extended the Julia sampler - which originally sampled the tensor-to-scalar ratio r and the lensing potential band power A_ϕ (Millea et al., 2020) - to jointly sample any subset of 5 of the Lambda Cold Dark Matter (LCDM) cosmological parameters, taking a prior value on the optical depth to reionization τ_{reio} due to its strong degeneracies with the other parameters. For map sizes up to 256^2 pixels the two implementations are comparable in run time; for larger maps some methods are a few factors slower in Python.

Statement of Need

We migrated **CMBLensing.jl** into Python for better integration with other Python-based machine learning and inference projects, planned and ongoing, that we are developing. Calling **CMBLensing.jl** from Python previously required the **juliacall** (JuliaPy, 2020) package and forced the user to write Julia inside a Python string, making debugging and iteration slow and tedious.

The LCDM parameter sampler is also a novel addition, expanding on the algorithm of (Millea et al., 2020) for jointly sampling r and A_ϕ . As far as we know, it is the first sampler to infer the LCDM parameters at the field level rather than from the CMB power spectra. We have used it to jointly infer $[\log(A_s), \theta_{MC}, \Omega_c h^2]$ from lensed temperature maps at varying white noise levels, with the other LCDM parameters fixed at fiducial values.

State of the Field

map_joint computes MAP estimates of the unlensed CMB fields and the lensing potential from a noisy lensed data map (or set of maps); here we compare its run time and accuracy between the two implementations. Both packages contain other CMB analysis tools, such as **lense_flow** (Millea et al., 2019) for gravitationally lensing CMB fields and **load_sim** for synthetic data generation, but **map_joint** is a good all around benchmark since it brings many methods together into a single tool. The LCDM sampler is left to the **Research Impact Statement** section, since it is novel and cannot be directly compared between languages.

38 We ran `map_joint` on temperature plus polarization data and stored the cached run time per
 39 language, averaged over 10 input data maps and 10 trials per map, at several square map sizes;
 40 the configuration is in Table 1, the results in Figure 1 and Table 2. Even after caching, run
 41 time varies between map seeds of the same size, due to the variable line search which finds the
 42 optimal step size at each iteration. The two implementations agree within 1 standard deviation
 43 at smaller map sizes, but at the largest map tested Python is up to 4 times slower than Julia.

Table 1: Configuration of the `map_joint` run time and accuracy comparison experiment. Every data map is generated by the Julia `load_sim` implementation and handed to both language implementations, so the two are timed on bit-for-bit identical inputs.

Quantity	Value
Data	Temperature plus polarization (IP)
Map sizes	128^2 , 256^2 , 512^2 , 1024^2 pixels
Resolution	2.5 arcminutes per pixel
White noise level	$3 \mu K$ arcminutes in T , $\sqrt{2} \times 3 \mu K$ arcminutes in E and B
$1/f$ noise	ℓ -knee = 100, α -knee = 3
Beam	None (FWHM = 0)
Mask	No pixel space sky mask. Fourier space low-pass held at unity through $\ell = 2950$, then a cosine ramp down to zero at $\ell = 3000$
<code>num_steps</code>	30
Data maps per map size	10
Trials per data map	10
Cached timing measurements per map size, per language	100




Figure 1: Run time comparison between language implementations of `map_joint`. The method is run twice in each language and the cached time is stored on the second call.

Table 2: Cached `map_joint` run times (mean $\pm$ standard deviation) versus map size for temperature plus polarization (IP) data. The same data as shown in Figure 1.

Map Size	Julia Average (s)	Python Average (s)	Python / Julia
128 ²	99.46 ± 41.7	192.9 ± 82.5	1.94
256 ²	550.4 ± 342	769.4 ± 300	1.4
512 ²	2332 ± 537	4799 ± 1440	2.06
1024 ²	6712 ± 1450	25700 ± 3930	3.83

Figures 2 and 3 compare accuracy. Figure 2 takes the data maps of Figure 1 and compares the average fractional difference between the Julia MAP estimates and the ground truth fields against the same quantity for Python. Only 4 series of lines are visible to the naked eye, so neither implementation is closer to the ground truth than the other; they agree to within numerical precision.




Figure 2: Fractional difference between the MAP estimates from either language and the simulation ground truth. Both use ground truth fields from the Julia simulation code, eliminating the possibility that the Python simulation and MAP code are wrong in a self-consistent way.

Figure 3 shows the average cross correlation of the ground truth fields with the Python and Julia MAP estimates, and of the two estimates with each other. The green lines stay nearly constant at 1, indicating the estimates agree on every length scale and are statistically identical.




Figure 3: Average cross correlation between Julia MAP x Julia Simulation, Python MAP x Julia Simulation, and Julia MAP x Python MAP for the 256^2 pixel maps of our comparison experiment. T , E , B , and ϕ are the temperature, E-mode, B-mode, and lensing potential fields. Blue and orange overlap completely, indicating statistically identical estimates.

Software Design

CMBLensing.py was migrated method by method, ensuring the two implementations produced outputs identical to within numerical precision for the same inputs. One of the first methods ported was `lense_flow`, which iteratively lenses a CMB temperature or polarization field given the unlensed fields and lensing potential (Millea et al., 2019). The Wiener filter, likelihood, and gradient methods followed, packaged into the main `map_joint` entry point.

Before implementing the new LCDM sampler, the machinery which jointly sampled (r, A_ϕ) was ported and confirmed to learn the same distributions between implementations for the same input maps. The step sampling those parameters was then swapped for one drawing LCDM parameter samples, leaving the neighboring field and lensing potential samplers untouched.

Unit tests in the `\tests` folder validate that corresponding methods in the two implementations produce identical outputs.

Research Impact Statement

Besides making CMBLensing natively available in JAX, the sampling algorithm in `sample_lcdm.py` represents the first field-level joint inference of the LCDM parameters that we are aware of. Other work, such as TOD0, has inferred them from the CMB power spectra, whereas we use image data as the input. The main entry point to the sampling algorithm is `sample_joint`. We include an example HPC slurm submission script in `\sampling_chains_TEMPLATE`.

During development we confirmed that each of the 5 sampled parameters could be learned in isolation with the other 4 held at ground truth. We also find that learning $[\log(A_s), \theta_{MC}, \Omega_c h^2]$

73 jointly is relatively easy for smaller temperature maps. The sampler could in theory be applied
 74 to larger subsets or the full set of 6 LCDM parameters, but in practice the map sizes needed
 75 to break the degeneracies and produce meaningful distributions with small widths become
 76 prohibitively large.

77 Table 3 outlines an experiment jointly inferring $[\log(A_s), \theta_{MC}, \Omega_c h^2]$ on 128×128 square
 78 pixel temperature maps.

Table 3: Configuration of the joint $[\log(A_s), \theta_{MC}, \Omega_c h^2]$ sampling experiment run with `sample_joint`.

Quantity	Value
Data	Temperature only (I)
Map size	128^2 pixels
Resolution	2.5 arcminutes per pixel
White noise level	$5 \mu K$ arcminutes
$1/f$ noise	None
Beam	None (FWHM = 0)
Mask	None
Sampled parameters	$\log(A_s), \theta_{MC}, \Omega_c h^2$
Parameters held at ground truth	$\Omega_b h^2, n_s, \tau_{\text{reio}}$
Ground truth	$[3.218387, 1.031732, 0.109381, 0.022386, 0.959814, 0.0]$
$[\log(A_s), \theta_{MC}, \Omega_c h^2, \Omega_b h^2, n_s, \tau_{\text{reio}}]$	
Search range, θ_{MC}	$[0.9328, 1.1452]$
Search range, $\log(A_s)$	$[2.661635, 3.782861]$
Search range, $\Omega_c h^2$	$[0.085, 0.155541]$
Chain starting points	Drawn at random inside the search ranges
θ step	Sequential 1 dimensional random walk
ϕ step	Metropolis, with proposal widths $\theta_{MC} = 3.5 \times 10^{-3}$, $\log(A_s) = 2.5 \times 10^{-2}$, $\Omega_c h^2 = 3 \times 10^{-3}$
ϕ initialization	HMC with <code>num_steps = 10</code> and <code>step_size = 0.05</code>
Iterations with θ held fixed	Zero
Data map realizations	100
Sub-chains per data map	50
Iterations per sub-chain	5
Burn-in discarded per sub-chain	6000
Thinning	600 samples
Convergence	By each sub-chain's integrated autocorrelation time
	Average Gelman-Rubin $\hat{R} < 1.05$

79 After running the 50 map realizations and 5 sub-chains per map of Table 3 for 6000 iterations
 80 each, burn-in was removed, the sub-chains thinned by their integrated auto-correlation time
 81 and concatenated into 50 processed chains, one per map. A Gaussian kernel density estimate
 82 on each gives 50 “per-map” distributions, whose product is taken and re-normalized. Its mean
 83 is our best estimate, across all input data maps, of a given parameter’s fiducial value. Figure 4
 84 shows the result for $\Omega_c h^2$; plots for $\log(A_s)$ and θ_{MC} are in the `\supplementary_materials`
 85 folder.




Figure 4: The resulting $\Omega_c h^2$ product distribution from 50 temperature data map realizations.

86 For the given noise level, map size, resolution, and sampler inputs we can compute a Fisher
 87 forecast for the expected lower bound on an individual map's learned uncertainty and compare
 88 it to the empirical widths. Figure 5 shows a triangle plot with Fisher in red and experimental
 89 results in blue. The experimental widths are only slightly larger, indicating our sampler performs
 90 close to optimally.

Posterior vs Fisher Forecast (Gaussian approximation, 1σ)




Figure 5: Triangle plot comparing the empirically found distributions and correlations with the Fisher forecast from the home-grown Fisher forecasting software in CMBLensing.py

Limitations

In theory our sampler could jointly infer all 6 LCDM parameters at once, but in practice the largest subset we could learn in testing was just 3. Many of the LCDM parameters are strongly degenerate, and breaking those degeneracies requires adding polarization data and increasing map size and resolution, all of which significantly slow down sampling. One iteration takes less than 3 seconds for 128×128 temperature-only maps with 4 cpus-per-task on Caltech's Resnick HPC (Caltech Resnick High Performance Computing Center, 2026), but preliminary testing suggests a joint inference of just 5 of the parameters (excluding τ_{reio}) could require temperature plus polarization maps of at least 1024×1024 square pixels, at minimum a $192\times$ increase in total run time. Running on GPU empirically gives at least a $4\times$ speed increase, which cannot counteract the $>192\times$ slowdown projected.

Besides these time complexity problems, researchers should be aware of the many tuning knobs which must be adjusted to get well-behaved, non-pathological chains. For a given map size, resolution, and noise level, a pilot chain must be run before any large HPC experiment to confirm optimal acceptance rates: around 65% for the ϕ HMC step (Beskos et al., 2010) and

around 44% for each of the 1-D Metropolis-Hastings steps of the N sampled LCDM parameters (Roberts et al., 1997). Tuning is time intensive: the user must wait for the samples to burn in, record the acceptance rates, and adjust the knobs, repeatedly, until the optimal rates are achieved.

Perhaps the most troublesome limitation is the current mixing scheme we employ to decouple the field and lensing potential in our sampler. We use the mixing matrices G and D defined in (Millea et al., 2020) which are computed in terms of the noise covariance matrix. For these matrix definitions, low SNR yields good mixing, but also large widths in the computed distributions for the most difficult parameters to learn, e.g. n_s or $\Omega_b h^2$. Learning these parameters would require a high SNR, but dropping the noise level diminishes the mixing matrices to the point that we either never burn in to the typical set or only ever sample the mode of the distribution. A mathematically justified mixing scheme decoupled from the experiment's noise would be a major area for future improvement, though we have been unable to find one.

AI Usage Disclosure

Claude Code was used for code generation in parts of the code base and for preparation of this manuscript. All AI generated code was reviewed by a human before being merged.

Acknowledgements

The computations presented here were conducted in the Resnick High Performance Computing Center, a facility supported by Resnick Sustainability Institute at the California Institute of Technology.

TODO who else do we wish to acknowledge?

References

- Beskos, A., Pillai, N. S., Roberts, G. O., Sanz-Serna, J. M., & Stuart, A. M. (2010). *Optimal tuning of the hybrid monte-carlo algorithm*. <https://arxiv.org/abs/1001.4460>
- Blood, Z. (2026). CMBLensing.py. In *GitHub repository*. GitHub. <https://github.com/zaneblood1/CMBLensing.py>
- Bradbury, J., Frostig, R., Hawkins, P., Johnson, M. J., Katariya, Y., Leary, C., Maclaurin, D., Necula, G., Paszke, A., VanderPlas, J., Wanderman-Milne, S., & Zhang, Q. (2018). *JAX: Composable transformations of Python+NumPy programs* (Version 0.3.13). <http://github.com/jax-ml/jax>
- Caltech Resnick High Performance Computing Center. (2026). *Caltech resnick high performance computing center*. <https://www.hpc.caltech.edu>
- JuliaPy. (2020). PythonCall.jl. In *GitHub repository*. GitHub. <https://github.com/JuliaPy/PythonCall.jl>
- Millea, M. (2017). CMBLensing.jl. In *GitHub repository*. GitHub. <https://github.com/marius311/CMBLensing.jl>
- Millea, M., Anderes, E., & Wandelt, B. D. (2019). Bayesian delensing of CMB temperature and polarization. *Physical Review D*, 100(2). <https://doi.org/10.1103/physrevd.100.023509>
- Millea, M., Anderes, E., & Wandelt, B. D. (2020). Sampling-based inference of the primordial CMB and gravitational lensing. *Physical Review D*, 102(12). <https://doi.org/10.1103/physrevd.102.123542>

148 Roberts, G. O., Gelman, A., & Gilks, W. R. (1997). Weak convergence and optimal scaling of
149 random walk Metropolis algorithms. *The Annals of Applied Probability*, 7(1), 110–120.
150 <https://doi.org/10.1214/aoap/1034625254>

DRAFT